

# Chapter4 Naive Bayes and Sentiment Classification

## 1、Text Classification

Sentiment analysis  
Spam detection  
Authorship identification  
Language Identification  
Assigning subject categories, topics, or genres

## 2、The Naive Bayes (NB) Classifier

1.

$$P(c | d) = \frac{P(d | c)P(c)}{P(d)}$$

2.

$$c_{MAP} = \operatorname{argmax}_{c \in C} P(c | d)$$

MAP is "maximum a posteriori" = most likely class

$$= \operatorname{argmax}_{c \in C} \frac{P(d | c)P(c)}{P(d)}$$

Bayes Rule

$$= \operatorname{argmax}_{c \in C} P(d | c)P(c)$$

Dropping the denominator

"Likelihood"

"Prior"

$$c_{MAP} = \operatorname{argmax}_{c \in C} P(d | c)P(c)$$

$$= \operatorname{argmax}_{c \in C} P(x_1, x_2, \dots, x_n | c)P(c)$$

Document d  
represented as  
features  
 $x_1 \dots x_n$

3.

"Likelihood"

"Prior"

$$c_{MAP} = \operatorname{argmax}_{c \in C} P(d | c) P(c)$$

$$= \operatorname{argmax}_{c \in C} P(x_1, x_2, \dots, x_n | c) P(c)$$

Document d  
represented as  
features  
x1..xn

4. 独立性假设

$$P(x_1, \dots, x_n | c) = P(x_1 | c) \cdot P(x_2 | c) \cdot P(x_3 | c) \cdot \dots \cdot P(x_n | c)$$

5.

$$c_{MAP} = \operatorname{argmax}_{c \in C} P(x_1, x_2, \dots, x_n | c) P(c)$$

$$c_{NB} = \operatorname{argmax}_{c \in C} P(c_j) \prod_{x \in X} P(x | c)$$

## NB Classifiers for Text Classification

1.

positions ← all word positions in test document

$$c_{NB} = \operatorname{argmax}_{c_j \in C} P(c_j) \prod_{i \in \text{positions}} P(x_i | c_j)$$

2. Problem with this formula

Multiplying lots of probabilities can result in floating-point underflow!

.0006 \* .0007 \* .0009 \* .01 \* .5 \* .000008....

Idea: Use logs, because  $\log(ab) = \log(a) + \log(b)$

We'll sum logs of probabilities instead of multiplying probabilities!

3. Solution

Instead of this:  $c_{NB} = \operatorname{argmax}_{c_j \in C} P(c_j) \prod_{i \in \text{positions}} P(x_i | c_j)$

This:  $c_{NB} = \operatorname{argmax}_{c_j \in C} \left[ \log P(c_j) + \sum_{i \in \text{positions}} \log P(x_i | c_j) \right]$

## NB: Learning

1.

$$\hat{P}(c_j) = \frac{N_{c_j}}{N_{total}}$$
$$\hat{P}(w_i | c_j) = \frac{count(w_i, c_j)}{\sum_{w \in V} count(w, c_j)}$$

2. Problem

$$\hat{P}(\text{"fantastic"} | \text{positive}) = \frac{count(\text{"fantastic"}, \text{positive})}{\sum_{w \in V} count(w, \text{positive})} = 0$$

3. Laplace (add-1) smoothing for Naïve Bayes

$$\hat{P}(w_i | c) = \frac{count(w_i, c) + 1}{\sum_{w \in V} (count(w, c) + 1)}$$
$$= \frac{count(w_i, c) + 1}{\left( \sum_{w \in V} count(w, c) \right) + |V|}$$

4. Training Processing

Calculate  $P(c_j)$  terms

- For each  $c_j$  in  $C$  do

$docs_j \leftarrow$  all docs with class  $= c_j$

$$P(c_j) \leftarrow \frac{|docs_j|}{|\text{total \# documents}|}$$

- Calculate  $P(w_k | c_j)$  terms

- $Text_j \leftarrow$  single doc containing all  $docs_j$

- For each word  $w_k$  in *Vocabulary*

$n_k \leftarrow$  # of occurrences of  $w_k$  in  $Text_j$

$$P(w_k | c_j) \leftarrow \frac{n_k + \alpha}{n + \alpha |Vocabulary|}$$

5. Unknown Word, Stop Words

**We ignore and remove them**

## Worked example

## A worked sentiment example with add-1 smoothing

	Cat	Documents
Training	-	just plain boring
	-	entirely predictable and lacks energy
	-	no surprises and very few laughs
	+	very powerful
	+	the most fun film of the summer
Test	?	predictable <del>with</del> no fun

### 1. Prior from training:

$$\hat{P}(c_j) = \frac{N_{c_j}}{N_{total}} \quad P(-) = 3/5 \\ P(+) = 2/5$$

### 2. Drop "with"

### 3. Likelihoods from training:

$$p(w_i|c) = \frac{\text{count}(w_i, c) + 1}{(\sum_{w \in V} \text{count}(w, c)) + |V|}$$

$$P(\text{"predictable"}|-) = \frac{1+1}{14+20} \quad P(\text{"predictable"}|+) = \frac{0+1}{9+20}$$

$$P(\text{"no"}|-) = \frac{1+1}{14+20} \quad P(\text{"no"}|+) = \frac{0+1}{9+20}$$

$$P(\text{"fun"}|-) = \frac{0+1}{14+20} \quad P(\text{"fun"}|+) = \frac{1+1}{9+20}$$

### 4. Scoring the test set:

$$P(-)P(S|-) = \frac{3}{5} \times \frac{2 \times 2 \times 1}{34^3} = 6.1 \times 10^{-5}$$

$$P(+ )P(S|+) = \frac{2}{5} \times \frac{1 \times 1 \times 2}{29^3} = 3.2 \times 10^{-5}$$

## Optimizing for Sentiment Analysis

Binary Multinomial Naive Bayes  
on a test document  $d$

First remove all duplicate words from  $d$

1. Then compute NB using the same equation:

$$c_{NB} = \operatorname{argmax}_{c_j \in C} P(c_j) \prod_{i \in \text{positions}} P(w_i | c_j)$$

2. dealing with negation

Simple baseline method:

Add NOT\_ to every word between negation and following punctuation:

didn't like this movie , but I



didn't NOT\_like NOT\_this NOT\_movie but I

3. Lexicon

## NB for other text classification tasks

1. Spam Filtering

提及数百万（美元）（（美元）NN，NNN，NNN.NN）

- 发件人：以许多数字开头
- 主题均为大写
- HTML的文本与图像区域的比率很低
- “百分之百保证”
- 您可以从列表中删除的索赔

## 2.NB in Language ID

确定文本是用什么语言编写的。基于字符n-grams的功能非常好。重要的是，要训练每种语言的多种多样性（例如，美国英语变体，如非洲裔美国英语，或世界各地的英语变体，例如印度英语）

## Summary

1. Very Fast, low storage requirements
2. Work well with very small amounts of training data
3. Robust to Irrelevant Features
4. Very good in domains with many equally important features
5. Optimal if the independence assumptions hold
6. A good dependable baseline for text classification

## 3.Evaluation : Pre, Rec, F1

		<i>gold standard labels</i>		
		gold positive	gold negative	
<i>system output labels</i>	system positive	<b>true positive</b>	<b>false positive</b>	<b>precision</b> = $\frac{tp}{tp+fp}$
	system negative	<b>false negative</b>	<b>true negative</b>	
		<b>recall</b> = $\frac{tp}{tp+fn}$		<b>accuracy</b> = $\frac{tp+tn}{tp+fp+tn+fn}$

## F-measure:

F measure: a single number that combines P and R:

$$F_{\beta} = \frac{(\beta^2 + 1)PR}{\beta^2 P + R}$$

We almost always use balanced  $F_1$  (i.e.,  $\beta = 1$ )

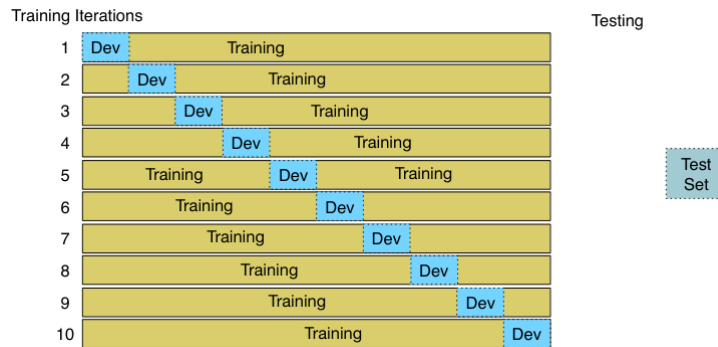
$$F_1 = \frac{2PR}{P + R}$$

Development Test Sets ("Devsets") and Cross-validation

在训练集上训练，在验证集上调优，在测试集上报告

- 这避免了过度拟合（“调整到测试集”）
- 对性能的更保守估计
- 但矛盾的是：需要尽可能多的数据用于训练，也需要尽可能多的数据用于测试；如何拆分

## Cross-validation: multiple splits



## Evaluation with more than two classes

		<i>gold labels</i>			
		urgent	normal	spam	
<i>system output</i>	urgent	8	10	1	$\text{precision}_u = \frac{8}{8+10+1}$
	normal	5	60	50	$\text{precision}_n = \frac{60}{5+60+50}$
	spam	3	30	200	$\text{precision}_s = \frac{200}{3+30+200}$
		$\text{recall}_u = \frac{8}{8+5+3}$	$\text{recall}_n = \frac{60}{10+60+30}$	$\text{recall}_s = \frac{200}{1+50+200}$	

Class 1: Urgent			Class 2: Normal			Class 3: Spam			Pooled		
	true urgent	true not		true normal	true not		true spam	true not		true yes	true no
system urgent	8	11	system normal	60	55	system spam	200	33	system yes	268	99
system not	8	340	system not	40	212	system not	51	83	system no	99	635
precision = $\frac{8}{8+11} = .42$			precision = $\frac{60}{60+55} = .52$			precision = $\frac{200}{200+33} = .86$			microaverage precision = $\frac{268}{268+99} = .73$		
macroaverage precision = $\frac{.42+.52+.86}{3} = .60$											

## 4.Statistical Significance Testing

### 统计显著性检验

统计显著性检验是一种通过检验观察数据与假设之间差异的统计方法。它的目的是确定一个样本数据是否有足够的证据支持假设，从而判断这个假设在整个总体中是否成立。

### 问题引入

鉴于：

- 分类器A和B
- 度量M：M(A, x) 是A在测试集x上的性能
- $\delta(x)$  : x上A、B之间的性能差异：
- $\delta(x) = M(A, x) - M(B, x)$
- 我们想知道 $\delta(x) > 0$ ，表示A优于B
- $\delta(x)$  称为效果大小
- 假设我们看到了 $\delta(x)$  为阳性。我们完成了吗？
- 不，这可能只是一个测试集的意外，或者是实验的环境。

## P-value

考虑两个假设

A不比B好  $H_0: \delta(x) \leq 0$

A比B好  $H_1: \delta(x) > 0$

想排除 $H_0$ ，我们在测试集上创建一个随机变量x

如果 $H_0$ 是真的，那么在这些测试集中，我们看到 $\delta(x)$ 的可能性有多大？

形式化为 p-value:

$P(\delta(x) \geq \delta(x) | H_0 \text{ 为真})$

P值是指观察数据出现在零假设 (null hypothesis) 下的概率。

具体地说：在零假设成立的情况下，得到观测结果或者比这个结果更极端的结果的概率。

如果P值非常小（通常设置一个置信水平 $\alpha(0.05, 0.01)$ 等），如果 $P < \alpha$ 我们拒绝原假设

## 如何计算P-value

分为参数性检测和非参数检测，在NLP中，我们不倾向于使用参数检验（如t检验），我们使用基于抽样的非参数测试

统计假设检验两种常见方法：

- 近似随机化
- 引导测试

Paired tests:

- 比较两组观测值，其中一组观测值中的每个观测值可以与另一组观测结果配对。
- 例如，当在同一测试集上查看系统A和B时，我们可以在每个相同的观察xi上比较系统A和系统B的性能

## The Paired Bootstrap Test

可以应用于任何度量（准确度、精度、召回率、F1等）。引导意味着从原始较大样本中重复抽取大量较小样本并替换（称为引导样本）

Calculate  $\delta(x)$  # how much better does algorithm A do than B on  $x$   
 $s = 0$

```

for  $j = 1$  to  $n$  do    # Draw a bootstrap sample  $x^{(i)}$  of size  $n$ 
    Select a member of  $x$  at random and add it to  $x^{(i)}$ 

```

$$s \leftarrow s + 1 \text{ if } \delta(x^{(i)}) > 2\delta(x)$$

```
return p-value(x)  # if very few did, our observed  $\delta$  is probably not accidental
```

Now we create, many, say,  $b=10,000$  virtual test sets  $x(i)$ , each of size  $n = 10$ .

[illegible]