

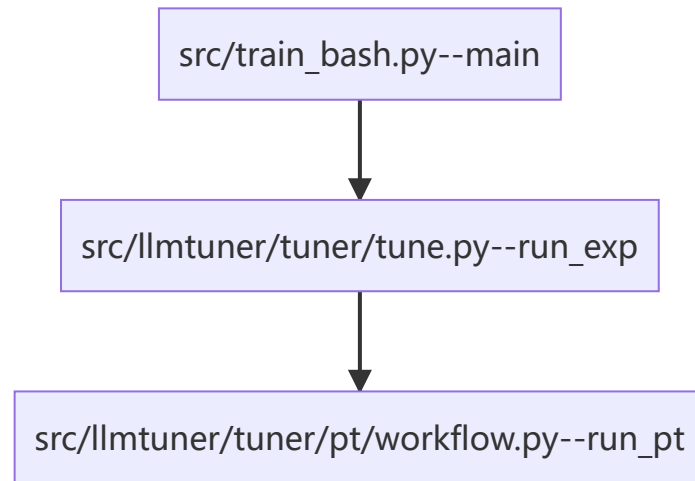
LLaMA-Efficient-Tuning

链接: <https://github.com/hiyouga/LLaMA-Efficient-Tuning>

```
1 LLaMA-Efficient-Tuning:
2 - assets: 微信群二维码
3 - data: 各种数据集 pt sft self-cognition multiturn rlhf
4 - src: 训练、推理代码
5     - llmtuner: 与微调相关的各种代码
6         - api: 创建FastAPI app的启动代码和数据对象定义
7         - chat: 利用模型进行问答, 返回形式包括流式和普通方式
8         - dsets: 数据集的校验, 加载, 预处理
9         - extras: 工具代码, 包含模型名字定义, 模板定义, 日志, 显存管理, 可视化等
10        - hparams: 阶段参数, 数据参数, 模型参数, 微调参数, 生成参数
11        - tuner: 核心代码, 包含dpo/ppo/pt/rm/sft
12        - webui: 模型微调与模型部署展示UI代码
13    - api_demo.py: 提供FastAPI接口
14    - cli_demo.py: 提供命令行接口
15    - export_model.py: 导出模型为单个或多个文件
16    - train_bash.py: 利用命令行命令训练模型
17    - train_web.py: 利用gradio web UI来训练模型
18    - web_demo.py: 利用gradio web UI来部署模型进行聊天
19 - tests: 测试代码
20     - evaluate_zh.py: 中文评测, 包含选择、填空、开放问答
21     - modeling_baichuan.py: 百川模型代码
22     - quantize.py: 将模型利用GPTQ算法量化至4bit
23     - template_encode.py: 测试模板编码
```

预训练代码解析

```
1 CUDA_VISIBLE_DEVICES=0 python src/train_bash.py \
2     --stage pt \
3     --model_name_or_path path_to_llama_model \
4     --do_train \
5     --dataset wiki_demo \
6     --template default \
7     --finetuning_type lora \
8     --lora_target q_proj,v_proj \
9     --output_dir path_to_pt_checkpoint \
10    --overwrite_cache \
11    --per_device_train_batch_size 4 \
12    --gradient_accumulation_steps 4 \
13    --lr_scheduler_type cosine \
14    --logging_steps 10 \
15    --save_steps 1000 \
16    --learning_rate 5e-5 \
17    --num_train_epochs 3.0 \
18    --plot_loss \
19    --fp16
```



```
1  #加载数据集
2  dataset = get_dataset(model_args, data_args)
3
4  #加载模型与分词器
5  model, tokenizer = load_model_and_tokenizer(model_args, finetuning_args,
6  training_args.do_train, stage="pt")
7
8  #处理数据集
9  dataset = preprocess_dataset(dataset, tokenizer, data_args, training_args,
10 stage="pt")
11
12 #获取数据收集器
13 data_collator = DataCollatorForLanguageModeling(tokenizer=tokenizer,
14 mlm=False)
15
16 # 初始化训练器
17 trainer = PeftTrainer(
18     finetuning_args=finetuning_args,
19     model=model,
20     args=training_args,
21     tokenizer=tokenizer,
22     data_collator=data_collator,
23     callbacks=callbacks,
24     **split_dataset(dataset, data_args, training_args)
25 )
26
27 # 训练
28 if training_args.do_train:
29     train_result =
30     trainer.train(resume_from_checkpoint=training_args.resume_from_checkpoint)
31     trainer.log_metrics("train", train_result.metrics)
32     trainer.save_metrics("train", train_result.metrics)
33     trainer.save_state()
34     trainer.save_model()
35     if trainer.is_world_process_zero() and model_args.plot_loss:
36         plot_loss(training_args.output_dir, keys=["loss", "eval_loss"])
37
38 # 评估, 使用困惑度
```

```

35 if training_args.do_eval:
36     metrics = trainer.evaluate(metric_key_prefix="eval")
37     try:
38         perplexity = math.exp(metrics["eval_loss"])
39     except OverflowError:
40         perplexity = float("inf")
41
42     metrics["perplexity"] = perplexity
43     trainer.log_metrics("eval", metrics)
44     trainer.save_metrics("eval", metrics)

```

GPT

```

1  class ResidualAttentionBlock(nn.Module):
2      def __init__(self, d_model: int, n_head: int, attn_mask: torch.Tensor =
None):
3          super().__init__()
4          #多头注意力层
5          self.attn = nn.MultiheadAttention(d_model, n_head)
6          #层规范化
7          self.ln_1 = LayerNorm(d_model)
8          #FFN层
9          self.mlp = nn.Sequential(OrderedDict([
10              ("c_fc", nn.Linear(d_model, d_model * 4)),
11              ("gelu", QuickGELU()),
12              ("c_proj", nn.Linear(d_model * 4, d_model))
13          ]))
14          #层规范化
15          self.ln_2 = LayerNorm(d_model)
16          #mask掉上三角的所有词语
17          self.attn_mask = attn_mask
18
19      def attention(self, x: torch.Tensor):
20          self.attn_mask = self.attn_mask.to(dtype=x.dtype, device=x.device)
21          if self.attn_mask is not None else None
22              return self.attn(x, x, x, need_weights=False,
attn_mask=self.attn_mask)[0]
23
24      def forward(self, x: torch.Tensor):
25          x = x + self.attention(self.ln_1(x))
26          x = x + self.mlp(self.ln_2(x))
27          return x
28
29  class Transformer(nn.Module):
30      def __init__(self, width: int, layers: int, heads: int, attn_mask:
torch.Tensor = None):
31          super().__init__()
32          self.width = width
33          self.layers = layers
34          self.resblocks = nn.Sequential(*[ResidualAttentionBlock(width,
heads, attn_mask) for _ in range(layers)])
35
36      def forward(self, x: torch.Tensor):
37          return self.resblocks(x)

```

```

37
38
39 def build_attention_mask(self):
40     # lazily create causal attention mask, with full attention between the
    vision tokens
41     # pytorch uses additive attention mask; fill with -inf
42     # 相当于将上三角（不包括对角线）的所有数值设为负无穷
43     mask = torch.empty(self.context_length, self.context_length)
44     mask.fill_(float("-inf"))
45     mask.triu_(1) # zero out the lower diagonal
46     return mask
47
48 self.transformer = Transformer(
49     width=transformer_width,
50     layers=transformer_layers,
51     heads=transformer_heads,
52     attn_mask=self.build_attention_mask()
53 )
54
55 def encode_text(self, text):
56     # 词语编码
57     x = self.token_embedding(text).type(self.dtype) # [batch_size, n_ctx,
    d_model]
58     # 位置编码加上词语编码得到输入
59     x = x + self.positional_embedding.type(self.dtype)
60     #将token长度维度放在第一维，Batchsize放在第二维
61     x = x.permute(1, 0, 2) # NLD -> LND
62     #输入Transformer得到输出
63     x = self.transformer(x)
64     #将token长度维度还原回第二维
65     x = x.permute(1, 0, 2) # LND -> NLD
66     #层规范化
67     x = self.ln_final(x).type(self.dtype)
68     # x.shape = [batch_size, n_ctx, transformer.width]
69     # take features from the eot embedding (eot_token is the highest number
    in each sequence)
70     x = x[torch.arange(x.shape[0]), text.argmax(dim=-1)] @
    self.text_projection
71
72     return x
73
74

```

permute理解

先将一个batch里同样位置的向量放在一起，这样直接使用第二维度的矩阵做乘法取对角线即为点积结果
最后再还原回来得到每个位置的输出向量

```
In [17]: x=torch.arange(24).reshape(2,3,4)
```

```
In [18]: x
```

```
Out[18]:
```

```
tensor([[[ 0,  1,  2,  3],
          [ 4,  5,  6,  7],
          [ 8,  9, 10, 11]],

        [[12, 13, 14, 15],
          [16, 17, 18, 19],
          [20, 21, 22, 23]]])
```

```
In [19]: x.permute(1,0,2)
```

```
Out[19]:
```

```
tensor([[[ 0,  1,  2,  3],
          [12, 13, 14, 15]],

        [[ 4,  5,  6,  7],
          [16, 17, 18, 19]],

        [[ 8,  9, 10, 11],
          [20, 21, 22, 23]]])
```